

Active Demo Selection for VLA Fine-Tuning

Project tracking deck — LIBERO/SmolVLA pipeline

anh0001 | branch `research/conformal-active-demo-budgeting`

updated 2026-05-28

One-line goal: collect fewer robot demos, train a better policy.

What we are trying to solve, in plain terms

Setting. A pretrained robot policy (SmolVLA) is great at *many* tasks out of the box, but it still has to be fine-tuned for a *new* task using human-teleoperated demonstrations.

The pain. Each demo is expensive — a person physically teleops the robot through the task. In sim we get demos for free, but the lesson should transfer to real Piper arms where demos cost minutes of human time.

The question.

- ▶ You have a pool of ~ 430 candidate demos (LIBERO-Spatial, 10 tasks).
- ▶ You have budget to train on **only 20**.
- ▶ **Which 20 should you pick?**

The hypothesis. Pick the demos the policy is *least sure* about $\Rightarrow$ same success rate at fewer demos.

The reality check. If active selection can't beat *random* on a strong pretrained model, the uncertainty signal is wrong — so far random is hard to beat, and that's the story this deck tracks.

What “input” and “output” actually mean here

INPUT

- ▶ A pretrained policy (SmolVLA-450M)
- ▶ A pool of ~ 430 unlabeled candidate episodes
- ▶ A budget schedule: $N = 5, 10, 15, 20$
- ▶ A query method (the thing we compare)


OUTPUT

- ▶ A *sample-efficiency curve*: success rate at each N .
- ▶ Plus a record of *which* episodes the method picked $\Rightarrow$ diagnostic for *why* it won or lost.

The query method is the only thing that changes between cells. Same backbone, same training recipe, same eval episodes. So if curves differ, the difference is the picking rule.

One “cell” = one (method, suite, seed). 4 training rounds per cell. ≈ 5 h wall-clock on one RTX 6000.

Pipeline at a glance



4 rounds per cell. One cell = (method, suite, seed). ≈5 h wall on one RTX 6000.

Query method 1: random

- ▶ **Input signal:** none.
- ▶ **Scoring rule:** `score = rng.random()` per candidate.
- ▶ **Selection rule:** top-5 by score (i.e. uniform pick of 5 from 30).
- ▶ **Expected failure mode:** bad demos by chance; high seed variance.
- ▶ **Why it matters:** the de-facto baseline; if a learner doesn't beat this on a well-pretrained VLA, the acquisition signal is anti-informative.

Query method 2: conformal [BROKEN]

- ▶ **Input signal:** teacher-forced policy loss per candidate episode (12 stride-8 frames, 1 forward pass each).
- ▶ **Scoring rule (as implemented):** per candidate,
 1. add its own loss to the calibration buffer,
 2. return $\text{loss}/\hat{q}$ if $\text{buffer} \geq 20$, else raw loss.
- ▶ **Selection rule:** top-5 by score.
- ▶ **Failure mode (confirmed):** self-referential calibration + per-candidate threshold + sorted candidate order $\Rightarrow$ structurally picks *highest-indexed* episodes in the random subsample, not the most uncertain.
- ▶ See “Conformal bug” slide for the full diagnosis.

Query method 3: entropy

- ▶ **Input signal:** same teacher-forced loss as conformal.
- ▶ **Scoring rule:** score = mean loss over sampled frames.
- ▶ **Selection rule:** top-5 by score.
- ▶ **Expected failure mode:** loss is dominated by nuisance (trajectory length, contact noise, gripper switches), not epistemic uncertainty. Picks redundant “hard” demos.
- ▶ **Empirical:** trails random by 6-12 pp at $N=20$ (1 seed). Consistent with the nuisance hypothesis.

Query method 4: dispersion

- ▶ **Input signal:** $K=4$ samples from SmolVLA's flow-matching action head per frame; mean per-(step, dim) std across samples.
- ▶ **Scoring rule:** score = mean dispersion over sampled frames.
- ▶ **Selection rule:** top-5 by score.
- ▶ **Why this should help:** measures self-disagreement in *action space*, not next-token likelihood. On a saved checkpoint: dispersion stayed in $[0.037, 0.043]$ while teacher-forced loss spanned $[0.04, 8.69]$ across 4 episodes — loss is mostly nuisance.
- ▶ **Failure mode (also confirmed empirically):** regresses at $N=15$ ($29 \rightarrow 22\%$). Without coverage, dispersion overselects from one visually-busy task.

Query method 5: `dispersion_quota`

- ▶ **Input signal:** same as `dispersion`.
- ▶ **Scoring rule:** same as `dispersion`.
- ▶ **Selection rule:** greedy round-robin under per-task coverage. For each pick, choose the candidate from the task with the smallest (budget + already-picked) count; break ties by score.
- ▶ **Why this should help:** forces task diversity; prevents the redundant-task collapse that pure `dispersion` shows.
- ▶ **Status (1 seed, 3 rounds):** holds $29 \rightarrow 31\%$ at $N=15$ where pure `dispersion` drops to 22% ; within ~ 2 pp of random at $N=15$.

Current results: libero_spatial, max_demos = 20

N	rand s0	rand s1	conf s0	conf s1	ent s0	disp s0	disp+q s0
5	9.0	10.5	9.0	10.5	9.0	9.0	9.0
10	22.5	31.0	21.0	16.0	15.5	29.0	29.0
15	33.0	32.5	24.5	18.5	25.5	22.0	31.0
20	35.0	40.5	28.0	25.0	29.0	OOM	OOM

- ▶ Eval is *paired* across methods (env seed = 1000+ run seed). Gaps are not eval noise.
- ▶ Grayed conformal columns: implementation is broken; kept for reproducibility, not direction-claim.
- ▶ **OOM**: GPU co-tenant returned mid- $N=20$ -train; partial curve only.

Conformal bug — how scoring went wrong

Codex's invariant check:

if conformal score = $\text{loss}/\hat{q}$ with a *global* $\hat{q}$, ranking = entropy ranking.

Empirical: seed-0 $N=5 \rightarrow 10$ new picks

- ▶ conformal: {1394, 1447, 1458, 1493, 1523}
- ▶ entropy: {1394, 1458, 1567, 1574, 1579}

Overlap 2/5. **Invariant violated.**

Smoking gun in `active_loop.py:77{82:`

- ▶ Each candidate's own loss is added to the buffer *before* it is scored (self-reference).
- ▶ Buffer resets every round; `calib_min=20`, `max_candidates=30`.
- ▶ Candidates 0–19 return raw loss (~ 0.03); 20–29 return $\text{loss}/\hat{q}$ (~ 1.0). $20 \times$ scale gap.
- ▶ Top-5 always lands in the last ~ 10 iterated candidates = highest-indexed in the sorted subsample.

Implication: conformal in this top- k pool setting reduces to a monotone transform; ranking $\equiv$ entropy. The branding adds no acquisition signal.

Status & next decision

Done

- ▶ Phase S.0–S.5 brought up. PS.4 single-task gate passed (80 % on libero_spatial task 0).
- ▶ 5 methods $\times$ 1–2 seeds each, full curves for 5 cells; 2 cells partial ($N = 5/10/15$) due to OOM.
- ▶ Conformal bug found and committed in the tracker (05549edf).

Blocked

- ▶ Both GPUs occupied by an unrelated user (HMDB51 ConvGRU, RTX 6000 $\times 2$).
- ▶ Pending: dispersion/dispersion_quota re-runs to fill the $N=20$ point, then seed-1 of each method.

Next decision after the ablation completes

- ▶ If dispersion_quota cleanly beats random at $N=15-20$ across seeds: pivot the paper to “diversity-aware action-space uncertainty,” bury conformal as a negative-result diagnostic.
- ▶ If quota only matches random: the publishable story is the *acquisition-signal pathology* itself (loss is anti-informative for low-demo VLA), plus the corrected recipe.